Aditya Shroff
101539669

PROGRAMMING FUNDAMENTALS FOR ANALYTICS

R_Assignment

Name: Aditya Shroff (101539669)
Professor: Mohsen Selseleh

Aditya Shroff
101539669

Step 1: Load the dataset and necessary libraries

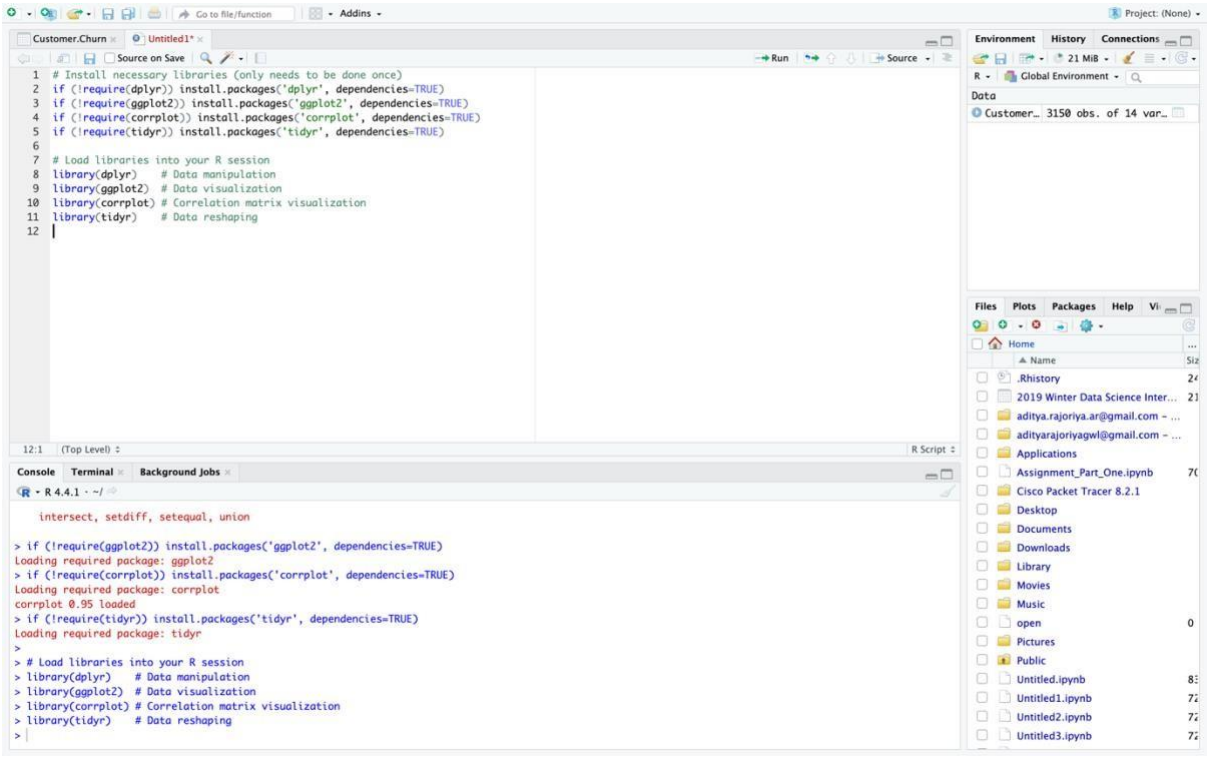

2. Load the Dataset and Handle Missing Values

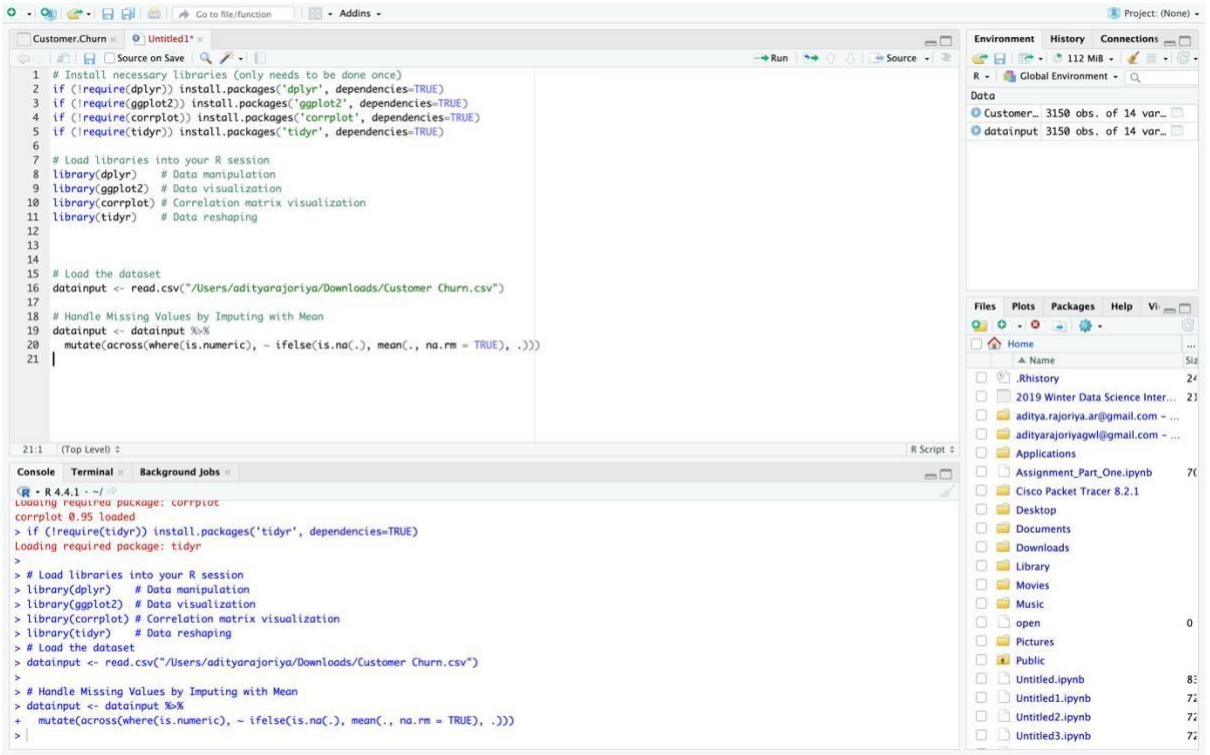

Aditya Shroff
101539669

3. Calculate IQR and Remove Outliers

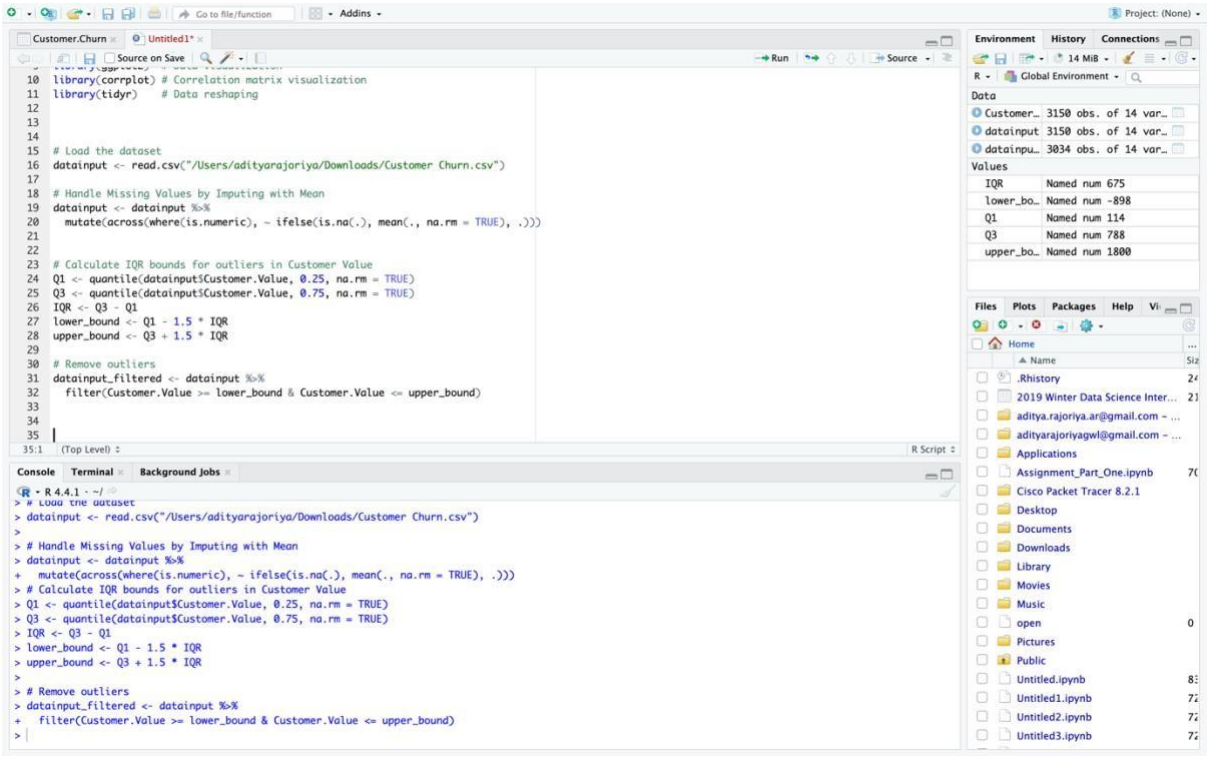

```
library(corrplot) # Correlation matrix visualization
library(tidyr)    # Data reshaping


# Load the dataset
datainput <- read.csv("/Users/adityarajoriya/Downloads/Customer Churn.csv")

# Handle Missing Values by Imputing with Mean
datainput <- datainput %>%
  mutate(across(where(is.numeric), ~ ifelse(is.na(.), mean(., na.rm = TRUE), .)))


# Calculate IQR bounds for outliers in Customer Value
Q1 <- quantile(datainput$Customer.Value, 0.25, na.rm = TRUE)
Q3 <- quantile(datainput$Customer.Value, 0.75, na.rm = TRUE)
IQR <- Q3 - Q1
lower_bound <- Q1 - 1.5 * IQR
upper_bound <- Q3 + 1.5 * IQR

# Remove outliers
datainput_filtered <- datainput %>%
  filter(Customer.Value >= lower_bound & Customer.Value <= upper_bound)
```

```
> # Load the dataset
> datainput <- read.csv("/Users/adityarajoriya/Downloads/Customer Churn.csv")
>
> # Handle Missing Values by Imputing with Mean
> datainput <- datainput %>%
+   mutate(across(where(is.numeric), ~ ifelse(is.na(.), mean(., na.rm = TRUE), .)))
> # Calculate IQR bounds for outliers in Customer Value
> Q1 <- quantile(datainput$Customer.Value, 0.25, na.rm = TRUE)
> Q3 <- quantile(datainput$Customer.Value, 0.75, na.rm = TRUE)
> IQR <- Q3 - Q1
> lower_bound <- Q1 - 1.5 * IQR
> upper_bound <- Q3 + 1.5 * IQR
>
> # Remove outliers
> datainput_filtered <- datainput %>%
+   filter(Customer.Value >= lower_bound & Customer.Value <= upper_bound)
>
```

4. Visualize Before and After Outlier Removal

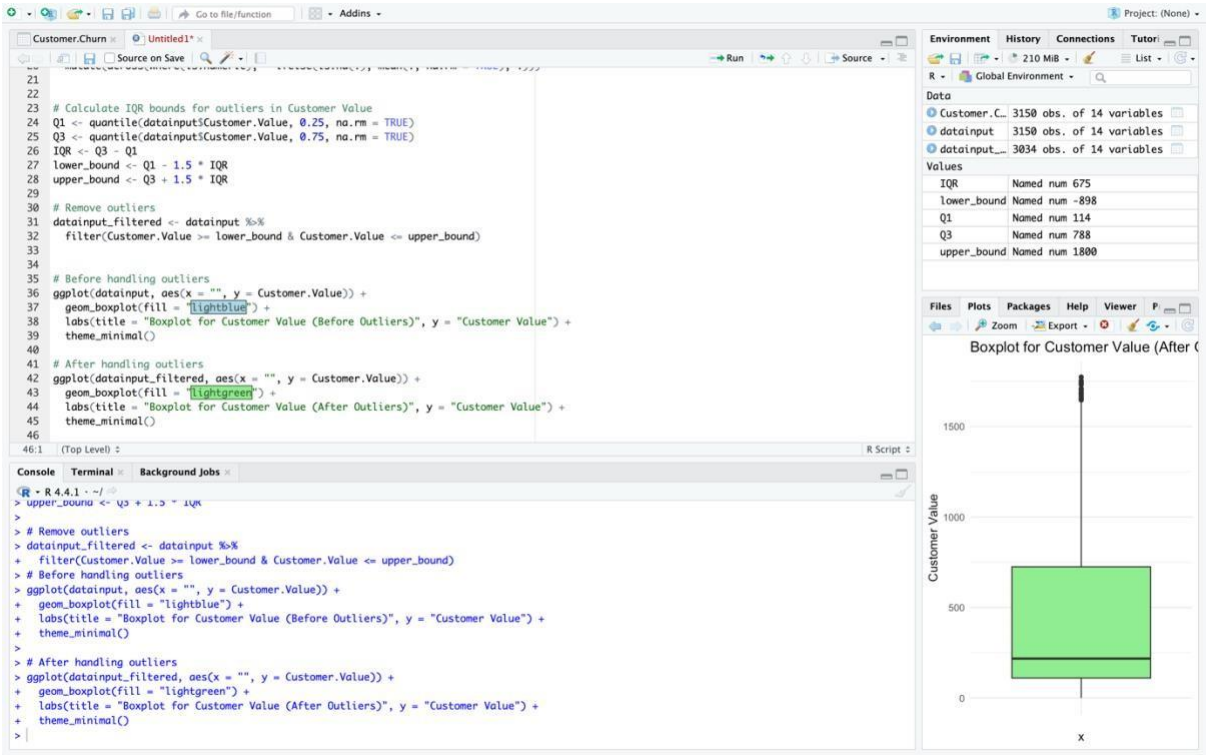

```
# Calculate IQR bounds for outliers in Customer Value
Q1 <- quantile(datainput$Customer.Value, 0.25, na.rm = TRUE)
Q3 <- quantile(datainput$Customer.Value, 0.75, na.rm = TRUE)
IQR <- Q3 - Q1
lower_bound <- Q1 - 1.5 * IQR
upper_bound <- Q3 + 1.5 * IQR

# Remove outliers
datainput_filtered <- datainput %>%
  filter(Customer.Value >= lower_bound & Customer.Value <= upper_bound)


# Before handling outliers
ggplot(datainput, aes(x = "", y = Customer.Value)) +
  geom_boxplot(fill = "lightblue") +
  labs(title = "Boxplot for Customer Value (Before Outliers)", y = "Customer Value") +
  theme_minimal()

# After handling outliers
ggplot(datainput_filtered, aes(x = "", y = Customer.Value)) +
  geom_boxplot(fill = "lightgreen") +
  labs(title = "Boxplot for Customer Value (After Outliers)", y = "Customer Value") +
  theme_minimal()
```

```
> # Remove outliers
> datainput_filtered <- datainput %>%
+   filter(Customer.Value >= lower_bound & Customer.Value <= upper_bound)
> # Before handling outliers
> ggplot(datainput, aes(x = "", y = Customer.Value)) +
+   geom_boxplot(fill = "lightblue") +
+   labs(title = "Boxplot for Customer Value (Before Outliers)", y = "Customer Value") +
+   theme_minimal()
>
> # After handling outliers
> ggplot(datainput_filtered, aes(x = "", y = Customer.Value)) +
+   geom_boxplot(fill = "lightgreen") +
+   labs(title = "Boxplot for Customer Value (After Outliers)", y = "Customer Value") +
+   theme_minimal()
>
```

Aditya Shroff
101539669

5. Group Data and Visualize Average Customer Value by Churn

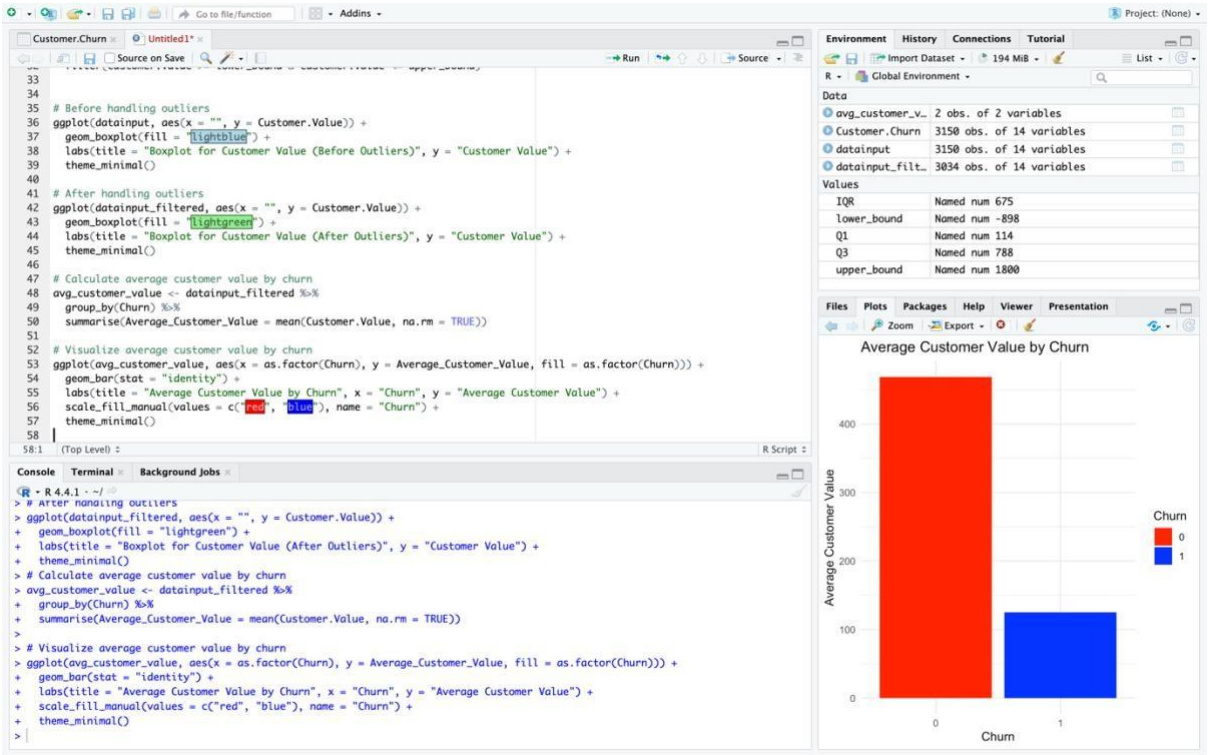

6. Create and Visualize Correlation Matrix

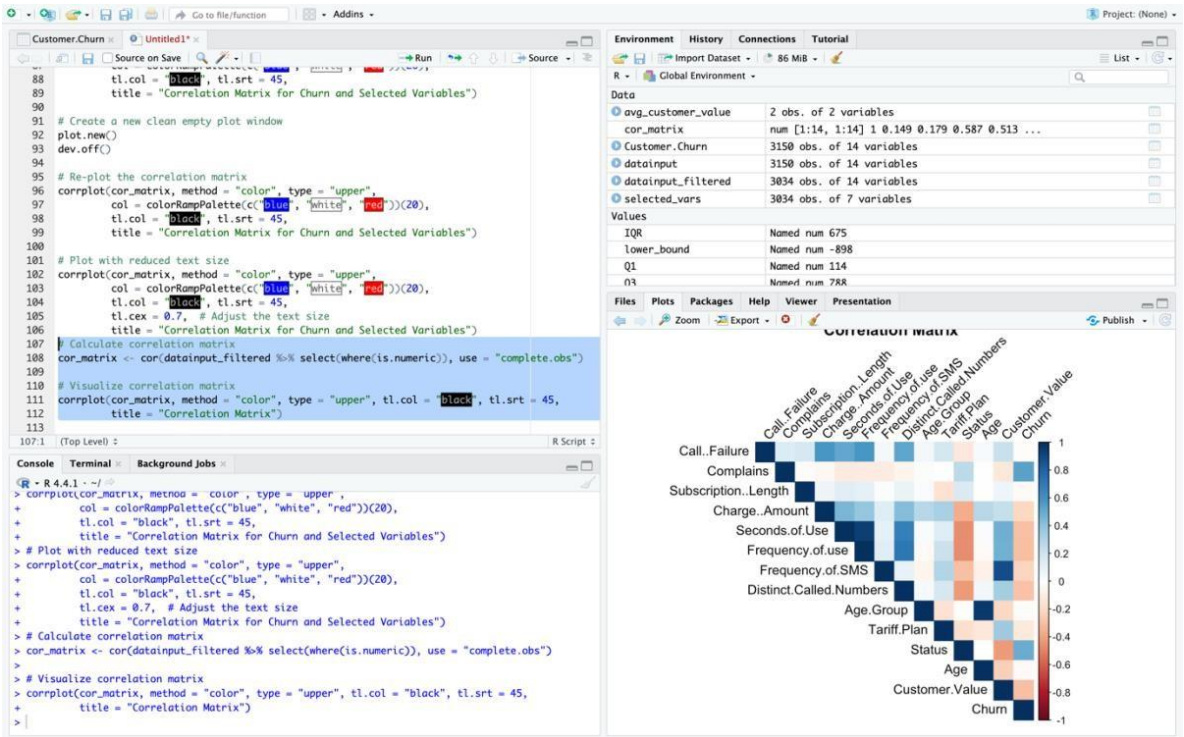

Aditya Shroff
101539669

7. Display Column Names

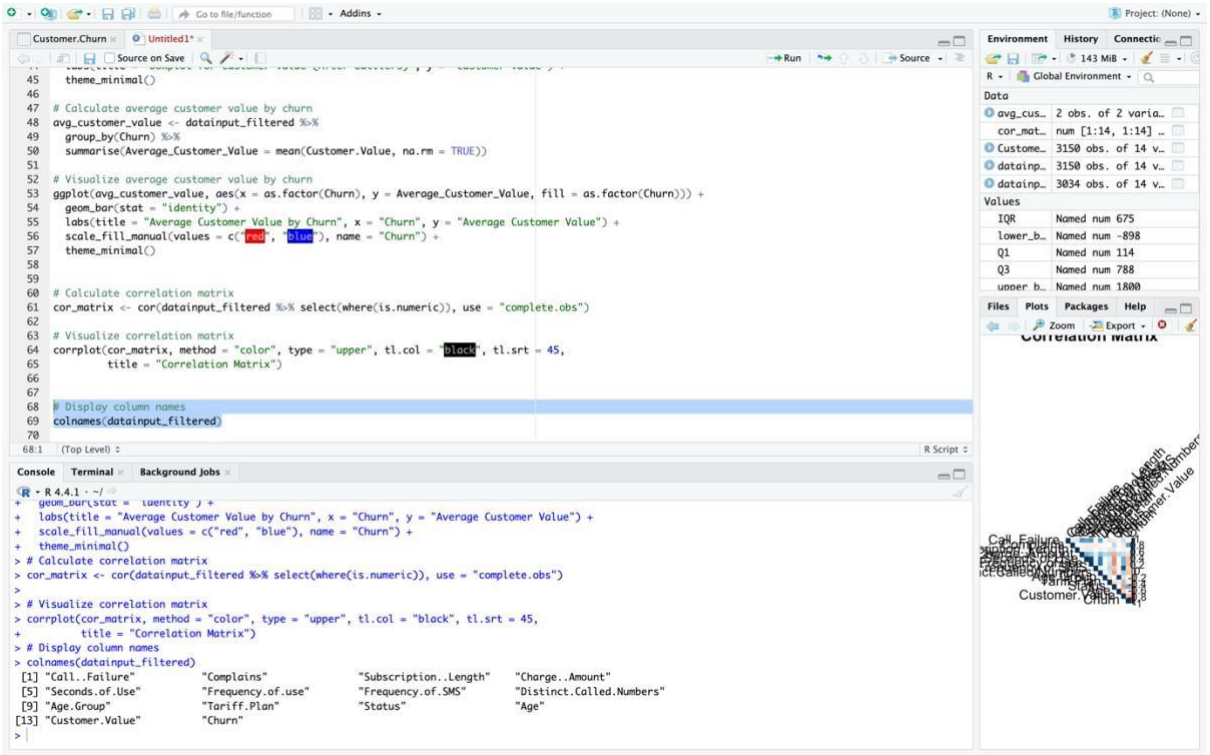

```
  theme_minimal()

# Calculate average customer value by churn
avg_customer_value <- datainput_filtered %>%
  group_by(Churn) %>%
  summarise(Average_Customer_Value = mean(Customer.Value, na.rm = TRUE))

# Visualize average customer value by churn
ggplot(avg_customer_value, aes(x = as.factor(Churn), y = Average_Customer_Value, fill = as.factor(Churn))) +
  geom_bar(stat = "identity") +
  labs(title = "Average Customer Value by Churn", x = "Churn", y = "Average Customer Value") +
  scale_fill_manual(values = c("red", "blue"), name = "Churn") +
  theme_minimal()


# Calculate correlation matrix
cor_matrix <- cor(datainput_filtered %>% select(where(is.numeric)), use = "complete.obs")

# Visualize correlation matrix
corrplot(cor_matrix, method = "color", type = "upper", tl.col = "black", tl.srt = 45,
         title = "Correlation Matrix")


# Display column names
colnames(datainput_filtered)
```

```
+   geom_bar(stat = "identity") +
+   labs(title = "Average Customer Value by Churn", x = "Churn", y = "Average Customer Value") +
+   scale_fill_manual(values = c("red", "blue"), name = "Churn") +
+   theme_minimal()
> # Calculate correlation matrix
> cor_matrix <- cor(datainput_filtered %>% select(where(is.numeric)), use = "complete.obs")
>
> # Visualize correlation matrix
> corrplot(cor_matrix, method = "color", type = "upper", tl.col = "black", tl.srt = 45,
+          title = "Correlation Matrix")
> # Display column names
> colnames(datainput_filtered)
 [1] "Call..Failure"          "Complains"              "Subscription..Length"   "Charge..Amount"
 [5] "Seconds.of.Use"         "Frequency.of.use"       "Frequency.of.SMS"       "Distinct.Called.Numbers"
 [9] "Age.Group"              "Tariff.Plan"            "Status"                 "Age"
[13] "Customer.Value"         "Churn"
>
```

8. Select Specific Variables and Create Another Correlation Matrix

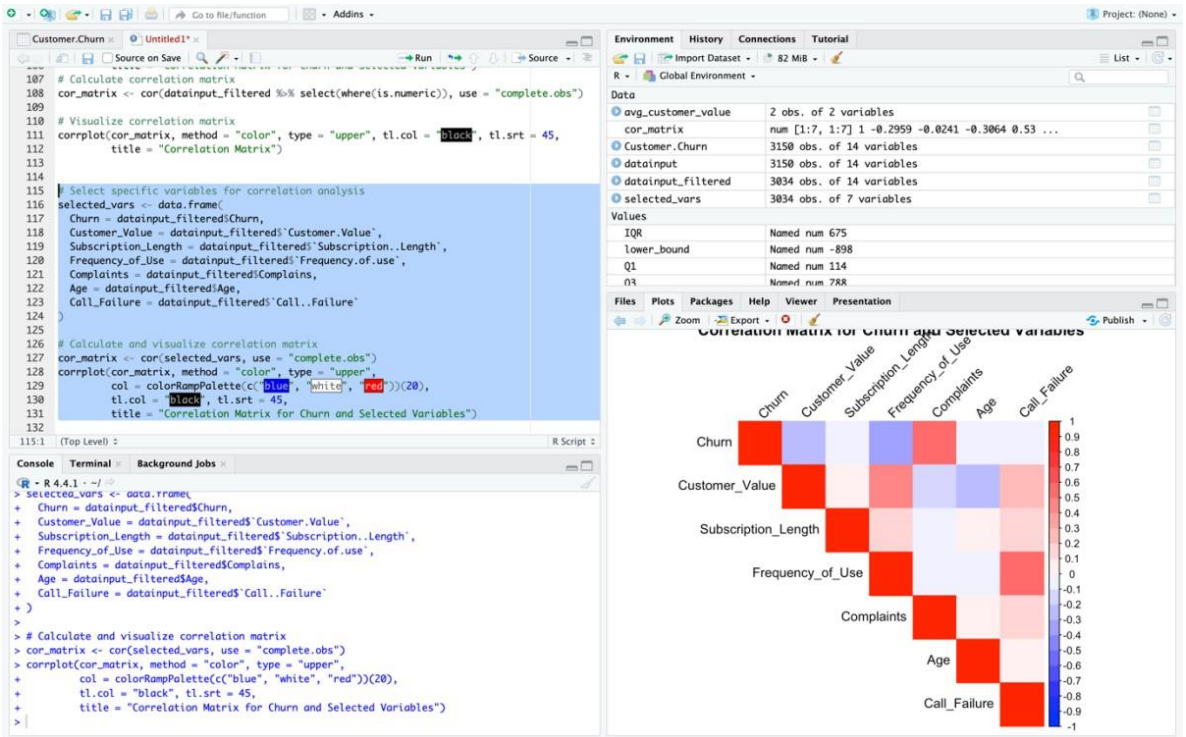

```
# Calculate correlation matrix
cor_matrix <- cor(datainput_filtered %>% select(where(is.numeric)), use = "complete.obs")

# Visualize correlation matrix
corrplot(cor_matrix, method = "color", type = "upper", tl.col = "black", tl.srt = 45,
         title = "Correlation Matrix")


# Select specific variables for correlation analysis
selected_vars <- data.frame(
  Churn = datainput_filtered$Churn,
  Customer_Value = datainput_filtered$`Customer.Value`,
  Subscription_Length = datainput_filtered$`Subscription..Length`,
  Frequency_of_Use = datainput_filtered$`Frequency.of.use`,
  Complaints = datainput_filtered$Complains,
  Age = datainput_filtered$Age,
  Call_Failure = datainput_filtered$`Call..Failure`
)

# Calculate and visualize correlation matrix
cor_matrix <- cor(selected_vars, use = "complete.obs")
corrplot(cor_matrix, method = "color", type = "upper",
         col = colorRampPalette(c("blue", "white", "red"))(20),
         tl.col = "black", tl.srt = 45,
         title = "Correlation Matrix for Churn and Selected Variables")
```

```
> selected_vars <- data.frame(
+   Churn = datainput_filtered$Churn,
+   Customer_Value = datainput_filtered$`Customer.Value`,
+   Subscription_Length = datainput_filtered$`Subscription..Length`,
+   Frequency_of_Use = datainput_filtered$`Frequency.of.use`,
+   Complaints = datainput_filtered$Complains,
+   Age = datainput_filtered$Age,
+   Call_Failure = datainput_filtered$`Call..Failure`
+ )
>
> # Calculate and visualize correlation matrix
> cor_matrix <- cor(selected_vars, use = "complete.obs")
> corrplot(cor_matrix, method = "color", type = "upper",
+          col = colorRampPalette(c("blue", "white", "red"))(20),
+          tl.col = "black", tl.srt = 45,
+          title = "Correlation Matrix for Churn and Selected Variables")
>
```

Aditya Shroff
101539669

9. Aggregate Data and Calculate Summary Statistics

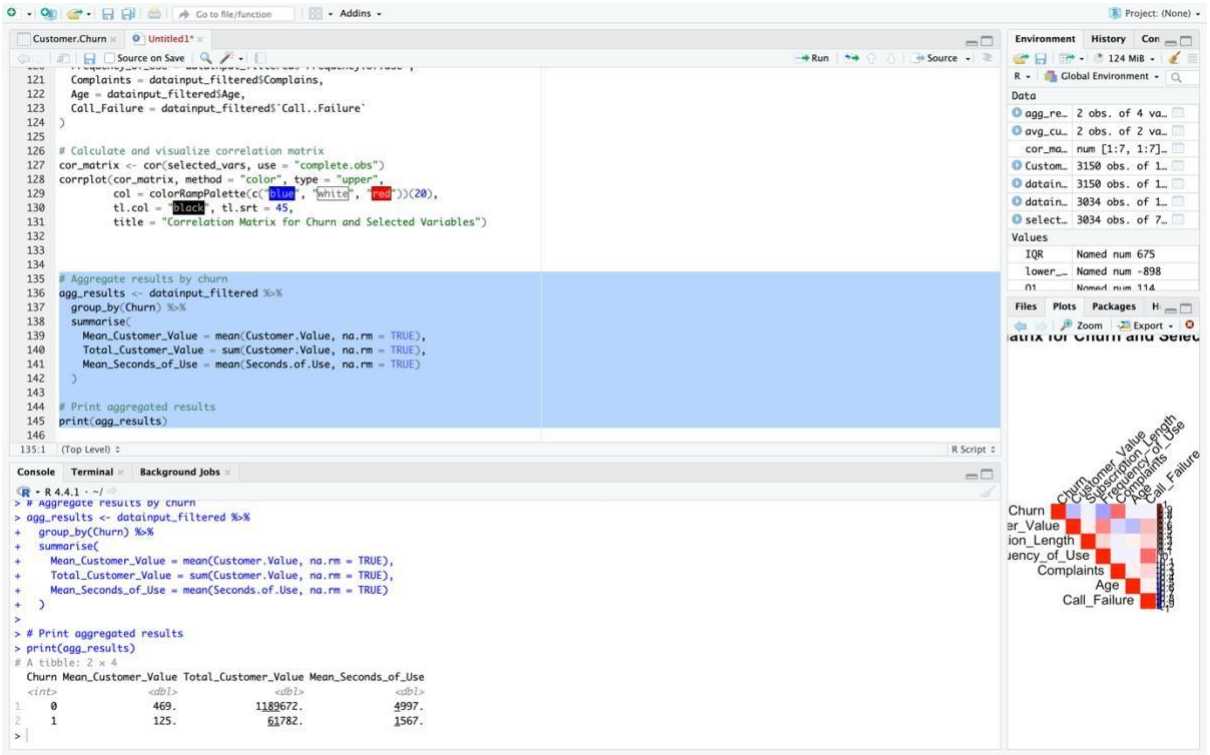

10. Perform T-Test

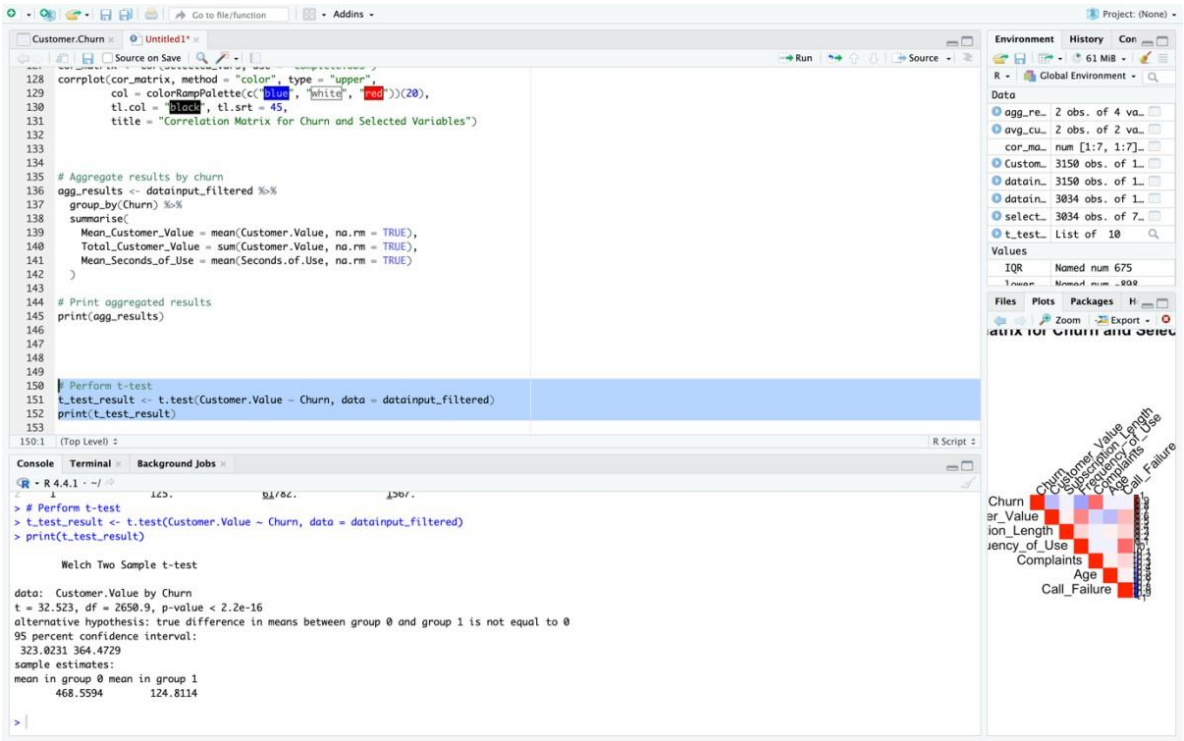

Aditya Shroff
101539669

11. Perform ANOVA

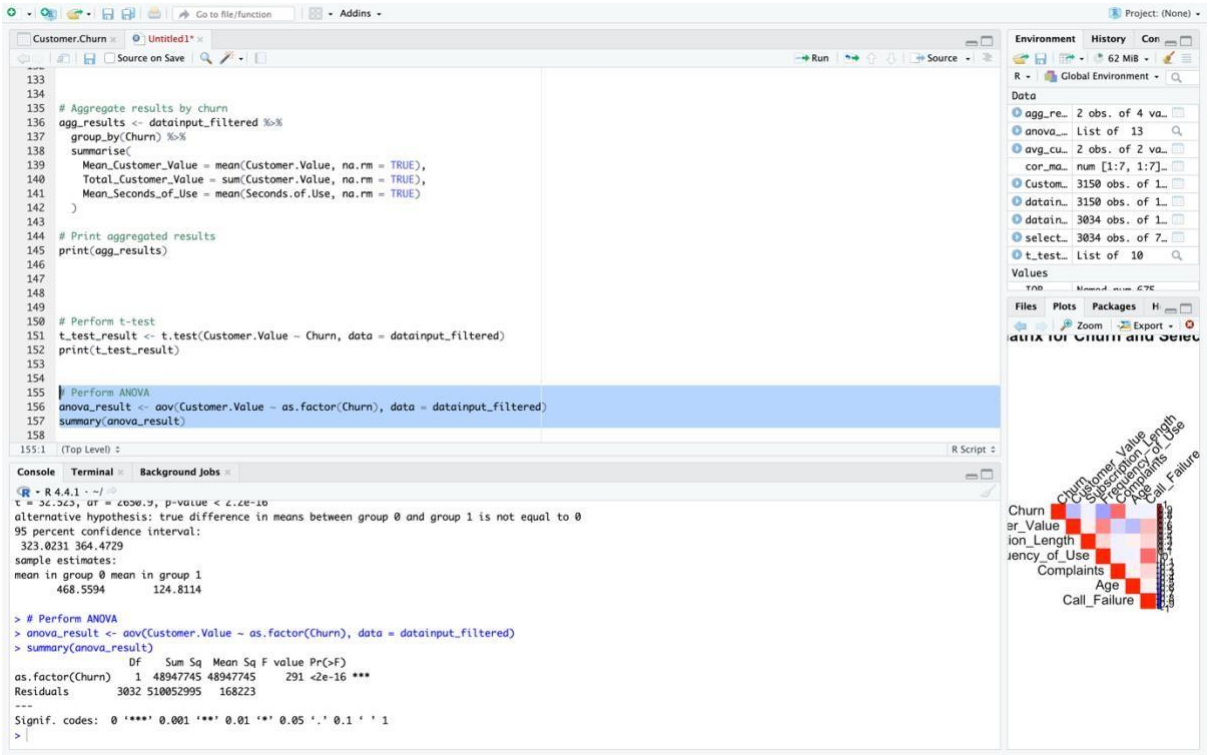

```
# Aggregate results by churn
agg_results <- datainput_filtered %>%
  group_by(Churn) %>%
  summarise(
    Mean_Customer_Value = mean(Customer.Value, na.rm = TRUE),
    Total_Customer_Value = sum(Customer.Value, na.rm = TRUE),
    Mean_Seconds_of_Use = mean(Seconds.of.Use, na.rm = TRUE)
  )

# Print aggregated results
print(agg_results)



# Perform t-test
t_test_result <- t.test(Customer.Value ~ Churn, data = datainput_filtered)
print(t_test_result)


# Perform ANOVA
anova_result <- aov(Customer.Value ~ as.factor(Churn), data = datainput_filtered)
summary(anova_result)
```

```
alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
95 percent confidence interval:
 323.0231 364.4729
sample estimates:
mean in group 0 mean in group 1
       468.5594        124.8114

> # Perform ANOVA
> anova_result <- aov(Customer.Value ~ as.factor(Churn), data = datainput_filtered)
> summary(anova_result)
                  Df    Sum Sq  Mean Sq F value Pr(>F)
as.factor(Churn)   1  48947745 48947745     291 <2e-16 ***
Residuals       3032 510052995   168223
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1
>
```

12. Build and Summarize Regression Model

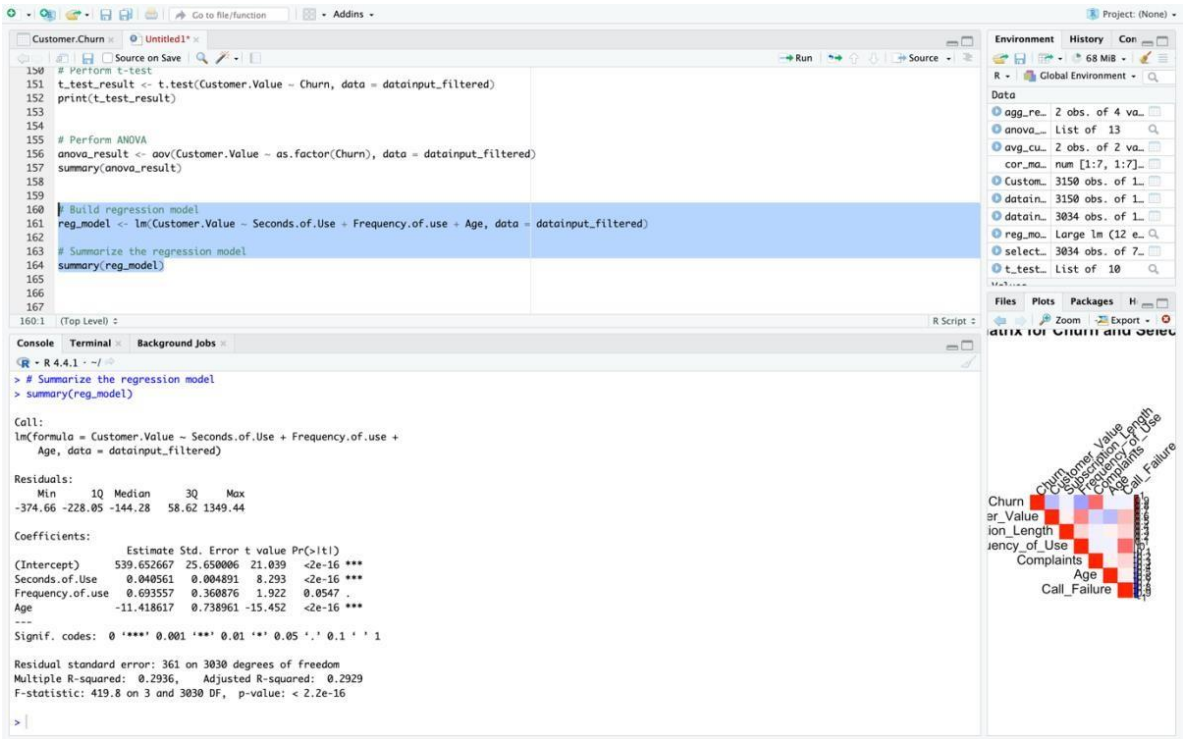

```
# Perform t-test
t_test_result <- t.test(Customer.Value ~ Churn, data = datainput_filtered)
print(t_test_result)


# Perform ANOVA
anova_result <- aov(Customer.Value ~ as.factor(Churn), data = datainput_filtered)
summary(anova_result)


# Build regression model
reg_model <- lm(Customer.Value ~ Seconds.of.Use + Frequency.of.use + Age, data = datainput_filtered)

# Summarize the regression model
summary(reg_model)
```

```
> # Summarize the regression model
> summary(reg_model)

Call:
lm(formula = Customer.Value ~ Seconds.of.Use + Frequency.of.use +
    Age, data = datainput_filtered)

Residuals:
    Min      1Q  Median      3Q     Max
-374.66 -228.05 -144.28   58.62 1349.44

Coefficients:
                Estimate Std. Error t value Pr(>|t|)
(Intercept)     539.652667  25.650006  21.039   <2e-16 ***
Seconds.of.Use    0.040561   0.004891   8.293   <2e-16 ***
Frequency.of.use  0.693557   0.360876   1.922   0.0547 .
Age             -11.418617   0.738961 -15.452   <2e-16 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 361 on 3030 degrees of freedom
Multiple R-squared:  0.2936,    Adjusted R-squared:  0.2929
F-statistic: 419.8 on 3 and 3030 DF,  p-value: < 2.2e-16

>
```